

Quiz MCQ – Final Assessment

Course: Full Stack AI B 4 SAT,SUN) 10 AM TO 12 PM

Dated:2nd Jan 2026

Name:

NIC No:

NumPy — 5 MCQs

1. What is the default data type of np.array([1, 2, 3])?

- A) float64
- B) int32 or int64 depending on platform
- C) object
- D) int8

1. Which operation broadcasts correctly?

```
1 import numpy as np
2 a = np.ones((3,1))
3 b = np.arange(3)
```

- A) a + b
- B) a + b.reshape(1,3)
- C) Both A and B
- D) Neither

1. What does np.ravel() return?

- A) A new flattened array (always a copy)
- B) A flattened view if possible, otherwise a copy
- C) A list
- D) A 2D array

1. What is the result of:

```
1 x = np.array([0, 1, 2, 3])
2 x[[True, False, True, False]]
```

- A) [0, 2]
- B) [1, 3]
- C) [0, 1, 2, 3]
- D) Error

1. Which creates a 2D array of shape (2, 3) filled with 7?

- A) np.full(2, 7, shape=(2,3))
 - B) np.full((2,3), 7)
 - C) np.fill((2,3), 7)
 - D) np.ones((2,3))*7.0
-

Python — Lists — 5 MCQs

1. What is the result of:

```
1 lst = [1, 2, 3]
2 lst.append([4,5])
3 lst
```

- A) [1,2,3,4,5]
 - B) [1,2,3,[4,5]]
 - C) Error
 - D) [1,2,3]
1. Which creates a shallow copy of a list?
- A) lst.copy()
 - B) lst[:]
 - C) list(lst)
 - D) All of the above

1. Result of:

```
1 a = [1,2]
2 b = a
3 b.append(3)
4 a
```

- A) [1,2]
 - B) [1,2,3]
 - C) Error
 - D) [3]
1. Which removes the first occurrence of value 2?
- A) lst.pop(2)
 - B) lst.remove(2)
 - C) del lst[2]

- D) lst.clear()
1. What does sorted([3, 1, 2]) return?
 - A) [3,2,1]
 - B) [1,2,3]
 - C) None
 - D) In-place sort
-

Python — Tuples — 5 MCQs

1. Which is a valid single-element tuple?
 - A) (1)
 - B) (1,)
 - C) tuple(1)
 - D) 1,()
1. What is true about tuples?
 - A) They are mutable
 - B) They are immutable
 - C) They cannot contain mutable items
 - D) They are ordered randomly
1. Result of:

```
1 t = (1, [2,3])
2 t[1].append(4)
3 t
```

- A) (1, [2,3])
 - B) (1, [2,3,4])
 - C) Error
 - D) (1, [4])
1. Which converts a list to a tuple?
 - A) tuple([1,2,3])
 - B) (1,2,3,)
 - C) list((1,2,3))

- D) tuple(1,2,3)
1. What does tuple unpacking require?
 - A) Same number of variables and tuple items
 - B) Only one variable
 - C) Names must match
 - D) Variables must be strings
-

Python — Sets — 5 MCQs

1. What is the result of set([1,2,2,3])?
 - A) {1,2,2,3}
 - B) {1,2,3}
 - C) [1,2,3]
 - D) Error
1. Which operation computes intersection?
 - A) a | b
 - B) a - b
 - C) a & b
 - D) a ^ b
1. Which structure is unhashable and cannot be added to a set?
 - A) int
 - B) str
 - C) tuple
 - D) list
1. Result of:

```
1 s = {1,2}
2 s.add(2)
3 s
```

- A) {1,2}
- B) {1,2,2}
- C) {2}
- D) Error

1. Symmetric difference of {1,2,3} and {3,4} is:

- A) {1,2,3,4}
 - B) {3}
 - C) {1,2,4}
 - D) {1,2}
-

Python — Dictionaries — 5 MCQs

1. Which retrieves a value safely with a default?

- A) d['key']
- B) d.get('key', default)
- C) d.pop('key', default)
- D) d.setdefault('key', default)

1. What happens in:

```
1 d = {'a':1, 'b':2}
2 d.update({'b':3, 'c':4})
```

- A) Error
- B) {'a':1, 'b':2, 'c':4}
- C) {'a':1, 'b':3, 'c':4}
- D) {'a':1}

1. Which iterates keys and values?

- A) for k in d:
- B) for k,v in d.items():
- C) for v in d.values():
- D) for k in d.keys():

1. What does dict.setdefault('x', []) do?

- A) Always overwrite 'x' with []
- B) Return value for 'x', create and set [] if missing
- C) Remove 'x'
- D) Throws KeyError if 'x' missing

1. Are dicts ordered?

- A) Never
 - B) Only in Python 3.7+ insertion order is preserved
 - C) Randomly ordered
 - D) Ordered by key
-

pandas — 5 MCQs

1. Which reads a CSV into a DataFrame?
 - A) pd.read_csv('file.csv')
 - B) pd.DataFrame('file.csv')
 - C) pd.load('file.csv')
 - D) pd.read('file.csv')
1. What does df.loc[2, 'col'] select?
 - A) Integer position 2 and column index 0
 - B) Label 2 for index and column 'col'
 - C) Slice rows up to 2
 - D) All rows, column 'col'
1. Which drops NA rows?
 - A) df.dropna()
 - B) df.fillna(0)
 - C) df.isna()
 - D) df.notna()
1. What does df.groupby('A')['B'].mean() return?
 - A) A single scalar
 - B) Series with mean of B per group in A
 - C) DataFrame with all columns
 - D) Error
1. How to create a new column as sum of two columns?
 - A) df['C'] = df.sum(axis=1)
 - B) df['C'] = df['A'] + df['B']
 - C) df['C'] = df.apply(sum)

- D) `df['C'] = df['A'].sum() + df['B'].sum()`
-

Seaborn — 5 MCQs

1. Which draws a scatterplot with hue by category?

```
1 import seaborn as sns
```

- A) `sns.lineplot(data=df, x='x', y='y', hue='cat')`
- B) `sns.scatterplot(data=df, x='x', y='y', hue='cat')`
- C) `sns.catplot(data=df, kind='bar', x='x', y='y')`
- D) `sns.histplot(data=df, x='x', hue='cat')`

1. Which function gives high-level facetting by a categorical variable?

- A) `sns.FacetGrid`
- B) `sns.PairGrid`
- C) `sns.catplot`
- D) `sns.clustermap`

1. Which palette argument is valid?

- A) `palette='viridis'`
- B) `palette='Set2'`
- C) `palette=['#1f77b4', '#ff7f0e']`
- D) All of the above

1. What does `sns.heatmap(df.corr(), annot=True)` do?

- A) Plots kernel density
- B) Plots correlation matrix with text annotations
- C) Plots pairwise scatter plots
- D) Plots bar chart

1. Which function draws distributions (univariate) with optional KDE?

- A) `sns.displot`
 - B) `sns.histplot`
 - C) `sns.kdeplot`
 - D) All can be used
-

scikit-learn — 10 MCQs

1. What does `train_test_split(X, y, test_size=0.2, random_state=42)` do?
 - A) Splits data; 20% test; reproducible split
 - B) Splits features only
 - C) Fits a model
 - D) Normalizes data

1. Which is a valid pipeline?

```
1 from sklearn.pipeline import Pipeline
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.linear_model import LogisticRegression
4 pipe = Pipeline([
5     ('scaler', StandardScaler()),
6     ('clf', LogisticRegression())
7 ])
```

- A) Valid; transforms then fit classifier
- B) Invalid; scalers cannot be used in pipeline
- C) Invalid order
- D) Pipelines require 3 steps

1. Which computes cross-validation scores?

- A) GridSearchCV
- B) `cross_val_score`
- C) KMeans
- D) PCA

1. In scikit-learn, how do you ensure reproducibility for random algorithms?

- A) Set `max_iter`
- B) Set `random_state`
- C) Set `n_jobs`
- D) Set `tol`

1. What does `OneHotEncoder(handle_unknown='ignore')` do?

- A) Raises error on unseen categories
- B) Drops all unknown rows
- C) Produces zeros for unknown categories

- D) Fills NaNs with mode
1. What is ColumnTransformer used for?
 - A) Combine multiple column-specific transformations
 - B) Select top features
 - C) Train multiple models
 - D) Plot correlations
 1. Which estimator is unsupervised?
 - A) RandomForestClassifier
 - B) LinearRegression
 - C) KMeans
 - D) SVC
 1. What does StandardScaler do?
 - A) Scales features to [0,1]
 - B) Centers to mean 0 and scales to unit variance
 - C) Normalizes rows to unit norm
 - D) Binarizes features
 1. Best way to avoid data leakage when preprocessing + model training?
 - A) Fit scaler on full dataset, then split
 - B) Use Pipeline so transforms fit on training only
 - C) Shuffle test labels
 - D) Drop NA values from test only
 1. Which metric is suitable for imbalanced binary classification?
 - A) Accuracy
 - B) MSE
 - C) ROC-AUC or F1-score
 - D) R²

TensorFlow / Keras — 10 MCQs

1. In Keras Sequential, the input_shape excludes:
 - A) Feature dimension

- B) Batch size
- C) Last layer size
- D) Number of classes

1. Correct binary classification compile?

```
1 model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])
```

- A) Correct
- B) Must use 'categorical_crossentropy'
- C) Must use 'mse'
- D) Metrics not allowed

1. What does model.fit(X, y, epochs=10, validation_split=0.2) do?

- A) Uses 20% of training data as validation
- B) Requires separate validation set
- C) Trains for 20 epochs
- D) Freezes layers

1. Which callback helps prevent overfitting?

- A) EarlyStopping
- B) ModelCheckpoint
- C) ReduceLROnPlateau
- D) All can help; EarlyStopping directly stops training

1. Output shape of:

```
1 model = tf.keras.Sequential([
2     tf.keras.layers.Dense(16, activation='relu', input_shape=(10,)),
3     tf.keras.layers.Dense(1, activation='sigmoid')
4 ])
```

- A) (None, 16)
- B) (None, 1)
- C) (None, 10)
- D) (1,)

1. In functional API, merging branches is done with:

- A) tf.concat only
- B) tf.keras.layers.concatenate()

- C) tf.add only
 - D) tf.keras.layers.Merge() (deprecated)
1. For multi-class classification with one-hot labels:
 - A) Use softmax + categorical_crossentropy
 - B) Use sigmoid + binary_crossentropy
 - C) Use linear + mse
 - D) Use relu + hinge
 1. How to freeze layers for transfer learning?
 - A) layer.trainable = False then re-compile
 - B) layer.active = False
 - C) model.freeze()
 - D) Not supported
 1. What does tf.data.Dataset.from_tensor_slices((X, y)).batch(32) create?
 - A) A generator for custom training loops only
 - B) A batched dataset usable with model.fit
 - C) A NumPy array
 - D) A Keras layer
 1. When using sparse_categorical_crossentropy, labels should be:
 - A) One-hot encoded
 - B) Integer class indices
 - C) Floats in [0,1]
 - D) Strings

LangChain / LLM — 5 MCQs

1. What does PromptTemplate do in LangChain?
 - A) Stores embeddings
 - B) Structure and format prompts with variables
 - C) Train LLMs
 - D) Vectorize documents
1. What are “Tools” in an Agent?

- A) Pretrained models
 - B) Functions/APIs the agent can call (e.g., search, calculator)
 - C) Memory modules
 - D) Datasets
1. What is a common way to enable retrieval-augmented generation (RAG)?
 - A) LLMChain only
 - B) Use a VectorStore + Retriever to fetch relevant docs for the prompt
 - C) Finetune the base LLM
 - D) Use TextSplitter alone
 1. What does ConversationBufferMemory do?
 - A) Stores embeddings in FAISS
 - B) Maintains chat history and injects it into the prompt
 - C) Schedules agent runs
 - D) Clears memory each turn
 1. Why use Callbacks (LangChain) during runs?
 - A) To manipulate optimizer
 - B) To observe/stream intermediate steps/tokens for logging or UX
 - C) To create databases
 - D) To cache prompts only